Agentforce Certification Notes – Set 1

Based on Randomized Question Screenshots (Q38–Q54)

Contents

1 Agents, Topics and Actions

1.1 How Agentforce Chooses Actions (Q43 & Q51)

- An **agent** is configured with:
  - One or more **topics** (use cases the agent can handle).
  - For each topic, a set of **actions** (flows, prompts, external services, etc.).
- When a user asks something:
  1. The **LLM interprets the user utterance** and identifies the most relevant topic.
  2. Within that topic, it chooses the **best matching actions** and the **correct order of execution**.
  3. The agent executes the selected actions and returns a grounded response.

Key Point

The **LLM's core role** in Agentforce is to understand user intent, map it to the right topic and actions, and decide their execution order. The configuration (topics + actions) constrains what it is allowed to do.

Exam Tip

If a question mentions *"How does Agentforce select the correct action"* or *"What is the role of the LLM in executing actions"*, the answer will emphasize:

- Matching the user utterance to topics and actions.
- Determining the order in which actions run.

1.2 Using Variables and Verification Steps (Q42)

- Security-related flows (like identity verification) often use **actions that output a status** (e.g., "verification_status = verified / not_verified").
- Best practice:
  1. Store the result of verification in a **variable** in the topic or flow.
  2. Apply **filters or conditions** so that other sensitive actions only execute when the variable confirms verification success.

Example

Use the output of a `Verify Identity` action to set an `isVerified` variable, then configure other service actions to run **only when** `isVerified = true`.

Exam Tip

If the question talks about *"ensuring verification is completed before any service-related actions"*, look for the option that:

- Stores the verification output in a variable, and
- Uses that variable as a condition or filter for subsequent actions.

Avoid answers that only rely on natural-language instructions like "Always verify identity" without variables or filters.

1.3 Custom Web Chat and Agent API (Q41)

- When customers must interact with an agent via a **custom website or custom chat UI**, you need a programmatic way to talk to the agent.
- **Agent API** provides this framework:
  - The web app sends messages and context to the agent.
  - The agent returns responses, actions, and state back to the web app.

Exam Tip

If the scenario mentions:

- "Custom web application",
- "Custom chat interface",
- "Framework for communication with the agent",

the correct choice is typically **Agent API**, not MCP or agent-to-agent integration.

2 Integrations and Channels

2.1 Slack Integration for Agents (Q53)

- To allow an Employee Agent or Service Agent to interact through **Slack**, you must first:
  - **Create a connection** between **Salesforce** and the target **Slack workspace**.
  - Then configure the appropriate Slack app / channel to route messages to the agent.

Exam Tip

Options mentioning Omni-Channel or embedded service deployments are more related to web chat / service channels, not directly to Slack workspaces. For Slack, think "**connect Salesforce and the Slack workspace**".

2.2 Digital Experience (Experience Cloud) + Messaging (Q46)

- When customers use an **Experience Cloud site** (Digital Experience) and you want **Messaging for In-App & Web**:
  1. Configure a **Messaging component** (Lightning component) on the Experience site page.

2. Connect that messaging channel to the Agentforce agent so logged-in users can chat.
3. Additional data like *membership number* can be passed through the messaging context.

Exam Tip

Look for answers involving:

- "Configure messaging Lightning component" or
- "Messaging for In-App and Web on the Digital Experience site."

Using MuleSoft or only Omni-Channel flows is not enough by itself for Messaging on a Digital Experience site.

2.3 Using Existing Backend Integrations (Q44)

- Often an enterprise already has integrations (e.g., to Oracle ERP) implemented via **flows** or other services.
- To leverage these with Agentforce:
  - Create a **custom agent action** that **calls the existing Flow**.
  - The LLM can then choose this action when the user asks about, for example, order fulfillment status.

Key Point

Don't rebuild integrations as pure prompt templates when a reliable backend flow already exists. Instead, **wrap the flow inside an agent action**.

3 Grounding, Data and Reporting

3.1 Grounding Data Quality (Q54)

- A prompt may be syntactically correct, but responses can still be poor if:
  - The **grounding data** is **incorrect, incomplete, or not relevant**.
  - The data source being retrieved by the agent is misconfigured.

Exam Tip

If the question says *"custom prompt summarizing case records is not returning appropriate summaries"*, the best explanation is usually that the **data used for grounding is incorrect or incomplete**, not that the Trust Layer or template version is wrong.

3.2 Agentforce Data Library and Knowledge Retrieval (Q38, Q40, Q45)

Purpose of the Data Library

- **Agentforce Data Library** is used to:
  - Store documents (policies, PDFs, knowledge articles, etc.).
  - Automatically **index** them for semantic search and retrieval.
  - Provide **grounded responses** based on those documents.

Example

UC stores HR policies, compliance guidelines, and procedures in a Data Library. The agent can then:

- Semantically search these documents,
- Ensure answers are grounded on published knowledge, and
- Reflect updates almost immediately as new versions are indexed.

External PDFs and Maintenance Guides (Q38)

- If PDF maintenance guides are stored externally (not in Salesforce) but must be used by a Service Agent:
  1. Upload the PDFs as **File** records into an **Agentforce Data Library**.
  2. Let the Data Library build a **search index**.
  3. Configure the agent to use this library as a **retrieval source** for answers.

Exam Tip

Copy-pasting external links into topic instructions is not enough. The exam will favor **Data Library + retrieval grounding** when the goal is robust, searchable documentation.

Using Past Cases as Knowledge (Q40)

- To answer questions based on similar, previously resolved cases:
  - Create an **Unstructured Data Model Object (UDMO)** or similar object based on the Case object.
  - Index case descriptions / resolutions for semantic similarity search.

Key Point

Past cases are **not** automatically used for reasoning. You must explicitly model and index them (often as unstructured text) for similarity-based retrieval.

3.3 Data Cloud DLOs for Agent Behavior (Q50)

- When analyzing agent behavior via Data Cloud, different **Data Lake Objects (DLOs)** represent different granularities:

  AI Agent Session High-level container capturing a **continuous interaction session** with one or more AI agents (e.g., a full chat from start to finish).

  AI Agent Interaction A specific interaction or step inside a session (e.g., a particular action call).

  AI Agent Interaction Message Individual messages within interactions (user + agent turns).

Exam Tip

If the question asks which DLO represents an *"overarching container"* or *"continuous interactions over time"*, the answer is **AI Agent Session**.

3.4 Enriched Event Logs (Q48)

- Enabling **"Enrich event logs with conversation detail"** in Agentforce configuration:
  - Stores **session-level detail**, including:
    - User input messages.
    - Agent responses for that session.
  - Helps in debugging, analytics, and compliance auditing.

Key Point

This setting makes event logs more verbose by embedding conversation content (not just metadata).

4 Prompt Builder, Templates and External Services

4.1 Sales Email Prompt Templates & Permissions (Q52)

- **Prompt templates** for Sales Emails are created in **Prompt Builder**.
- To allow the **sales team** to use those templates:
  1. Enable **Sales Emails** in Setup.
  2. Assign the appropriate **Prompt Template permission set** (e.g., *Prompt Template User* for consumers; *Prompt Template Manager* for creators and admins).

Exam Tip

Data Cloud Admin permissions do not control prompt template usage. Focus on:

- Sales Email feature toggle in Setup, and
- Prompt Template-specific permission sets.

4.2 Prompt Flows and External Services (Q47)

- A **template-triggered prompt flow** may call an **external service** (REST / gRPC) to get data.
- To use returned data inside a prompt template:
  - Surface the response as **External Service Record merge fields**.
  - Reference these merge fields in the prompt template so the LLM can see the external data before generation.

Example

An external service returns the current account balance. The flow maps that response into an External Service Record, which then appears as a merge field like `{{ExternalService.Balance}}` in the prompt template.

4.3 Model Context Protocol (MCP) Use Cases (Q49)

- **Model Context Protocol (MCP)** lets the LLM:
  - Dynamically discover and use tools/APIs at runtime.
  - Connect to external systems without hard-coding every endpoint in advance.

Example

A **legal assistant agent** uses MCP to:

1. Discover a document classification API at runtime.
2. Call that API to analyze case files.
3. Use the classification results as context for its legal analysis.

Exam Tip

If a scenario talks about an agent *discovering and using external tools/APIs dynamically*, that is a strong signal for **MCP**. Do not confuse this with:

- Agent Cards (describing capabilities of other agents), or
- Simple real-time collaboration between two agents.

5 Testing and Environments

5.1 Agentforce Testing Center Best Practices (Q39)

- **Agentforce Testing Center** allows you to run test cases (often from CSV) to validate agent behavior.
- To avoid modifying real CRM data:
  - Run tests **in a sandbox environment**, not in production.
  - Use representative but non-sensitive data in test cases.

Exam Tip

If the question asks how to avoid data changes while testing, the correct answer is typically:

Use the Testing Center only in a sandbox environment.

Limiting the number of test cases or running tests in production does not solve data safety concerns.

6 Quick Revision Checklist (Q38–Q54)

- **Q54** – Wrong summaries ⇒ check **grounding data quality**, not just the Trust Layer.
- **Q53** – Employee Agent + Slack ⇒ **create Salesforce–Slack workspace connection**.
- **Q52** – Sales Email prompt templates ⇒ **Sales Emails enabled + Prompt Template permission set**.

- **Q51 & Q43** – LLM chooses **topic + actions + order** based on user utterance.
- **Q50** – Overarching container DLO ⇒ **AI Agent Session**.
- **Q49** – Dynamic API discovery ⇒ **Model Context Protocol (MCP)**.
- **Q48** – Enrich event logs ⇒ **store user input + agent responses** in session logs.
- **Q47** – Use external service data in prompts ⇒ **External Service Record merge fields**.
- **Q46** – Agent on Experience Cloud site ⇒ **Messaging Lightning component** hooked to agent.
- **Q45** – Live policy & compliance knowledge ⇒ **Agentforce Data Library** as retrieval source.
- **Q44** – Oracle ERP fulfillment status ⇒ **custom agent action calling an existing Flow**.
- **Q42** – Identity verification status ⇒ store in **variable** and apply **filters/conditions**.
- **Q41** – Custom website chat ⇒ **Agent API**.
- **Q40** – Similar past cases ⇒ create **(unstructured) data model object** and index it.
- **Q39** – Testing Center safety ⇒ run in **sandbox**, not production.
- **Q38** – External PDFs as knowledge ⇒ upload to **Agentforce Data Library** and use retrieval.