

MuleSoft Design Patterns

Integration Architecture Advanced Flows



Integration
Mastery

Scalability · Reliability · Reusability

Topics Covered:

API-Led Connectivity · Batch Processing
Scatter-Gather · System Sync · Large Files

Prepared for:

Senior MuleSoft Developer / Architect
December 9, 2025

Contents

1 API-Led Connectivity Pattern	2
1.1 Definition & Purpose	2
1.2 Architecture Diagram	3
1.3 Best Practices	3
2 Scatter-Gather (Parallel Processing)	4
2.1 Definition	4
2.2 Internal Working	4
2.3 Architecture Diagram	4
2.4 Handling the Output (DataWeave)	4
2.5 Common Mistakes	5
3 Batch Processing (ETL)	6
3.1 When to use it?	6
3.2 Internal Architecture (Phases)	6
3.3 Mule XML Configuration	6
4 System Synchronization & Watermarking	8
4.1 The Concept	8
4.2 Watermarking (Object Store)	8
4.2.1 Flow Logic	8
4.3 Scheduler	8
5 Large File Processing	9
5.1 The Challenge: Memory	9
5.2 Streaming Strategies	9
5.2.1 1. Repeatable File Store Stream	9
5.3 Using For-Each vs. Batch	9
6 Interview Guide	10
6.1 Critical Interview Questions	10
6.1.1 Q1: How does Scatter-Gather handle variables?	10
6.1.2 Q2: What is the difference between For-Each and Batch?	10
6.1.3 Q3: How do you handle a "Composite Routing Error" in Scatter-Gather?	10
6.2 Mini Project: Sync Service	10

Chapter 1

API-Led Connectivity Pattern

1.1 Definition & Purpose

API-Led connectivity is an architectural approach that shifts away from point-to-point integration to a structured, reusable, three-layered architecture.

Purpose:

- **Decoupling:** Changes in the backend (e.g., swapping SAP for Oracle) do not break the frontend apps.
- **Reusability:** A "Customer API" can be reused by Mobile, Web, and B2B systems.
- **Agility:** Different teams can work on different layers simultaneously.

1.2 Architecture Diagram

Experience Layer (XAPI)

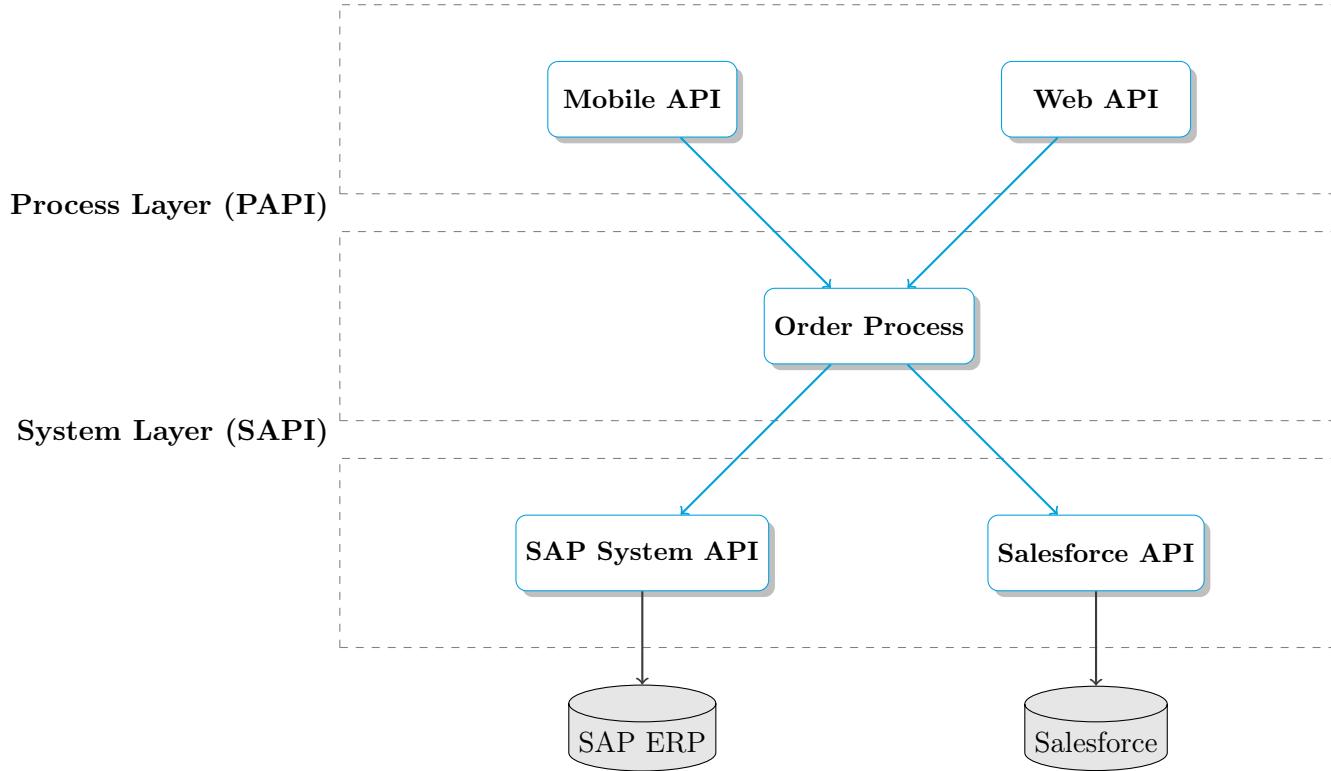


Figure 1.1: The Three Layers of API-Led Connectivity

1.3 Best Practices

- **XAPI:** Should contain NO business logic. Only format conversion (JSON/XML).
- **PAPI:** Where the logic lives (aggregation, orchestration, routing).
- **SAPI:** Should mirror the backend system. Simple CRUD operations.

Chapter 2

Scatter-Gather (Parallel Processing)

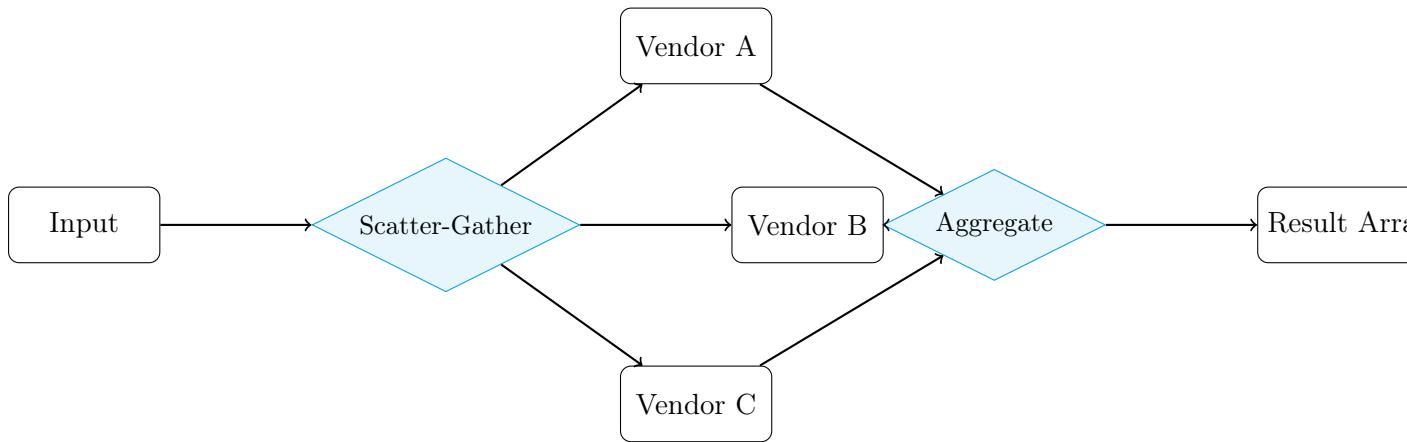
2.1 Definition

The **Scatter-Gather** router sends a message to multiple routes *concurrently* (in parallel) and waits for all of them to complete. It then aggregates the results into a single payload.

2.2 Internal Working

1. **Multicasting:** The event is cloned (deep copy) and sent to each route.
2. **Threading:** Each route runs in a separate thread (CPU-Light or I/O-Blocking pools).
3. **Aggregation:** The output is a Mule Message Object containing a collection of results.

2.3 Architecture Diagram



2.4 Handling the Output (DataWeave)

The output of Scatter-Gather is an Object where keys are numeric indices (0, 1, 2...). We often need to flatten this.

```
1 %dw 2.0
2 output application/json
3 ---
4 // This extracts just the payloads from the multipart structure
5 flatten(payload..payload)
```

Listing 2.1: Processing Scatter-Gather Output

2.5 Common Mistakes

Failure Behavior

If **one** route fails, the entire Scatter-Gather fails and throws a MULE:COMPOSITE_ROUTING error.

Solution: Wrap *each* route inside a **Try Scope** with an **On Error Continue**. This ensures that if Vendor A fails, Vendor B and C still return data, and you handle the partial failure in the aggregation step.

Chapter 3

Batch Processing (ETL)

3.1 When to use it?

Use Batch Processing when dealing with:

- Large datasets (e.g., 100,000+ records).
- Data synchronization (ETL).
- Reliability is key (handling individual record failures).

3.2 Internal Architecture (Phases)

1. **Load Dispatch:** The payload is split into individual records and stored in a persistent queue (serialization).
2. **Process (Steps):** Records are processed asynchronously in blocks (default block size: 100).
3. **On Complete:** A summary report is generated (Total records, Success count, Failure count).

3.3 Mule XML Configuration

```
1 <batch:job name="syncContactsBatch" maxFailedRecords="-1">
2
3     <batch:process-records>
4         <batch:step name="ValidateStep">
5             <ee:transform>
6                 <ee:message>
7                     <ee:set-payload><![CDATA [%dw 2.0
8                         output application/java
9                             -->
10                            payload update { case .email -> lower(.) }
11                            ]]></ee:set-payload>
12                     </ee:message>
13                 </ee:transform>
14             </batch:step>
15
16             <batch:step name="InsertStep" accept-policy="ONLY_FAILURES">
17                 </batch:step>
18
19             <batch:step name="BulkInsertStep">
```

```
20      <batch:aggregator doc:name="Batch Aggregator" size="100">
21          <db:insert doc:name="Bulk Insert" ... />
22      </batch:aggregator>
23  </batch:step>
24 </batch:process-records>
25
26 <batch:on-complete>
27     <logger message="Batch Finished: #[payload]"/>
28 </batch:on-complete>
29
30 </batch:job>
```

Listing 3.1: Batch Job Structure

Batch Aggregator

Always use **Batch Aggregator** for Database or Salesforce inserts. Inserting 100 records in 1 API call is significantly faster than 100 separate API calls.

Chapter 4

System Synchronization & Watermarking

4.1 The Concept

Synchronizing data between two systems (e.g., pulling new employees from HR to Active Directory) requires keeping track of what has already been processed.

4.2 Watermarking (Object Store)

Watermarking is the technique of storing the timestamp or ID of the last processed record.

4.2.1 Flow Logic

1. **Retrieve:** Get ‘lastModifiedDate‘ from Object Store. (Default to 1900-01-01 if null).
2. **Query:** `SELECT * FROM Users WHERE ModifiedDate > ‘vars.watermark‘.`
3. **Process:** Sync the records.
4. **Store:** Update the Object Store with the ‘max(ModifiedDate)‘ from the current batch.

4.3 Scheduler

Synchronization is usually triggered by a **Scheduler** component (e.g., run every 15 minutes).

Fixed Frequency vs Cron

- **Fixed Frequency:** Runs every X minutes from the *end* of the last run.
- **Cron:** Runs at specific wall-clock times (e.g., every Friday at 5 PM). Beware of overlapping runs if the process takes too long!

Chapter 5

Large File Processing

5.1 The Challenge: Memory

Processing a 2GB CSV file in a standard JVM (which might only have 1GB heap) will cause an ‘OutOfMemoryError’.

5.2 Streaming Strategies

Mule 4 uses streaming by default, but for large files, specific configurations ensure stability.

5.2.1 1. Repeatable File Store Stream

Mule keeps a small buffer in memory and writes the rest to a temporary file on disk. This allows you to read a stream multiple times without crashing RAM.

```
1 <file:config name="File_Config">
2   <file:connection workingDir="/tmp/input" />
3 </file:config>
4
5 <ee:transform>
6   <ee:message>
7     <ee:set-payload><![CDATA [%dw 2.0
8       output application/json deferred=true
9       ---
10      payload map (item) -> { ... }
11    ]]></ee:set-payload>
12   </ee:message>
13 </ee:transform>
```

Listing 5.1: Streaming Config

5.3 Using For-Each vs. Batch

For large files (CSV/XML):

- **Batch:** Best for record-level error handling and ETL. Overhead of serialization.
- **For-Each:** Lighter. Use with `batchSize` (e.g., split 1000 records at a time) to avoid loading everything.

Chapter 6

Interview Guide

6.1 Critical Interview Questions

6.1.1 Q1: How does Scatter-Gather handle variables?

Answer: Variables set *inside* a route are **local** to that route. They are NOT propagated to the aggregation phase or other routes. Only the payload and attributes are aggregated.

6.1.2 Q2: What is the difference between For-Each and Batch?

Answer:

- **For-Each:** Synchronous, single-threaded. If one fails, the loop stops (unless Try-Catch used). No persistent queuing.
- **Batch:** Asynchronous, multi-threaded. Persistent queues. Handles failures gracefully (continues to next record). Includes reporting phase.

6.1.3 Q3: How do you handle a "Composite Routing Error" in Scatter-Gather?

Answer: This error occurs when one route fails. To access the partial success data, checking the error object is required, but the best practice is to wrap every route in a Try Scope with On-Error-Continue, returning a specific structure (e.g., 'success: false, error: ...') so the aggregator can filter them out manually.

6.2 Mini Project: Sync Service

Scenario: Sync Contacts from Database to Salesforce every 10 mins.

1. **Source:** Scheduler (10 min).
2. **Logic:** Retrieve Watermark -> Select DB (WHERE date > watermark).
3. **Processing:** Use **Batch Job**.
4. **Batch Step 1:** Transform DB structure to Salesforce JSON.
5. **Batch Step 2 (Aggregator):** Use Salesforce Create-Bulk connector (Commit size 200).
6. **On Complete:** Update Watermark in Object Store.